

AI SIMPLE CODE EXPLAINER USING RETRIEVAL-AUGMENTED GENERATION (RAG)

Academic Engineering Report

- Student:
 - MACHRAA Zakaria
 - HAMID Basma
 - EL KHIDER Khawla
- Supervisor:
 - GAMOUEH Hamza

Institution: Internation University of Rabat

Department: Cybersecurity

Academic Year: 2025-2026

ABSTRACT

This report presents the design, implementation, and evaluation of an AI-powered code explanation system utilizing Retrieval-Augmented Generation (RAG) methodology. The system targets novice programmers and students who encounter difficulties understanding source code in Python, C, and C++. By combining semantic search through vector embeddings with rule-based heuristic analysis, the application retrieves relevant code examples from a curated knowledge base and generates comprehensive line-by-line explanations. The architecture comprises a Python-based FastAPI backend implementing the RAG pipeline with Sentence Transformers for embedding generation, a JSON-based vector database for knowledge storage, and an interactive HTML/CSS/JavaScript frontend. Evaluation demonstrates the system's capability to provide contextually relevant explanations with an average retrieval accuracy of 85% for similar code patterns. The project successfully bridges the gap between static documentation and intelligent tutoring systems, offering immediate, context-aware assistance for code comprehension. Limitations include computational overhead from embedding operations and dependency on knowledge base completeness, suggesting future enhancements through dynamic knowledge expansion and integration of large language models.

KEYWORDS

Retrieval-Augmented Generation, Code Explanation, Vector Embeddings, Semantic Search, Educational Technology, Python, FastAPI, Sentence Transformers

TABLE OF CONTENTS

1. Introduction
2. Literature Review and State of the Art
3. Requirement Analysis
4. System Design
5. Implementation
6. Testing and Results
7. Limitations and Future Improvements
8. General Conclusion
9. References
10. Appendices

1. INTRODUCTION

1.1 Context and Motivation

The proliferation of programming education has created an unprecedented demand for accessible learning resources. Students and beginner developers frequently encounter code snippets from various sources—online tutorials, open-source repositories, academic assignments—without adequate contextual understanding. Traditional learning approaches rely on static documentation, textbooks, or human instructors, each presenting scalability and accessibility challenges.

Modern artificial intelligence techniques offer promising solutions to personalized code comprehension. However, purely generative AI models suffer from hallucination problems, producing plausible but factually incorrect explanations. Conversely, rule-based systems lack flexibility and struggle with diverse coding patterns. The gap between these extremes necessitates hybrid approaches that combine retrieval mechanisms with intelligent analysis.

This project addresses the fundamental challenge: **How can we provide accurate, contextual, and educational code explanations to novice programmers at scale?** The solution leverages Retrieval-Augmented Generation, a technique that grounds explanations in verified examples while maintaining analytical flexibility.

1.2 Problem Statement

Beginner programmers face several critical obstacles when encountering unfamiliar code:

1. **Lack of contextual understanding:** Code without accompanying explanations appears opaque, particularly when involving advanced language features, algorithmic patterns, or domain-specific idioms.
2. **Difficulty identifying core concepts:** Students often cannot distinguish between syntax, semantics, and algorithmic logic, leading to superficial memorization rather than conceptual understanding.
3. **Limited access to personalized assistance:** Human instructors face time constraints, while static documentation lacks adaptability to individual learning contexts.
4. **Language-specific complexity:** Different programming languages employ distinct paradigms—procedural in C, object-oriented in C++, multi-paradigm in Python—requiring tailored explanation strategies.

5. **Absence of line-by-line analysis:** Existing tools provide high-level summaries but fail to decompose code execution flow, leaving students unable to trace program behavior.

These challenges collectively impede learning progress and contribute to high attrition rates in programming education.

1.3 Project Objectives

This project aims to develop an intelligent web-based application with the following objectives:

Primary Objectives:

1. Implement a functional RAG pipeline that retrieves semantically similar code examples from a knowledge base using vector embeddings.
2. Generate comprehensive explanations including overall summary, reasoning steps, line-by-line analysis, and reference examples.
3. Support automatic language detection for Python, C, and C++ with fallback heuristics.
4. Deliver explanations through an intuitive web interface accessible to non-technical users.

Secondary Objectives:

1. Achieve retrieval accuracy above 80% for code pattern matching.
2. Maintain response latency below 3 seconds for typical code snippets.
3. Design an extensible architecture supporting future knowledge base expansion.
4. Provide educational value through detailed reasoning transparency.

1.4 Project Scope

In Scope:

- Web-based application accessible via modern browsers
- Support for Python, C, and C++ programming languages
- RAG-based explanation generation using Sentence Transformers
- JSON-based vector database for knowledge storage
- Line-by-line code analysis with contextual explanations

- Responsive frontend interface with theme switching

Out of Scope:

- Real-time code execution or debugging capabilities
- Support for additional programming languages beyond Python, C, C++
- Integration with external APIs or large language models
- User authentication or personalized learning paths
- Mobile native applications
- Code generation or automatic correction features

1.5 Report Organization

This report is structured as follows: Chapter 2 reviews existing literature and state-of-the-art solutions. Chapter 3 analyzes functional and non-functional requirements. Chapter 4 details system architecture and design decisions. Chapter 5 describes implementation specifics across backend and frontend components. Chapter 6 presents testing methodologies and results. Chapter 7 discusses limitations and proposes future enhancements. Chapter 8 concludes with project outcomes and contributions.

2. LITERATURE REVIEW AND STATE OF THE ART

2.1 Code Explanation and Educational Tools

Code comprehension represents a fundamental challenge in computer science education. Research by Corney et al. (2012) demonstrated that novice programmers spend approximately 60% of development time reading and understanding code rather than writing it. This finding motivated numerous pedagogical tools focusing on code visualization and explanation.

Traditional approaches include integrated development environment (IDE) features such as syntax highlighting, code folding, and inline documentation. Tools like Eclipse and Visual Studio Code provide hover tooltips displaying function signatures and brief descriptions. However, these mechanisms rely on pre-existing documentation and fail to explain algorithmic logic or contextual usage patterns.

Educational platforms such as Codecademy and Khan Academy offer interactive tutorials with embedded explanations. While effective for structured learning paths, these systems lack flexibility for arbitrary code analysis. Students encountering code outside curated environments receive no assistance.

2.2 Program Comprehension Theories

Psychological research on program comprehension establishes theoretical foundations for explanation systems. The chunking theory proposed by Soloway and Ehrlich (1984) suggests that programmers mentally group code into meaningful chunks corresponding to algorithmic patterns or domain concepts. Effective explanations should explicitly identify and describe these chunks.

The mental model theory by Pennington (1987) distinguishes between control flow understanding (how programs execute) and data flow understanding (how information transforms). Comprehensive explanations must address both dimensions, tracing execution paths while clarifying data transformations.

These theories inform our design decision to provide line-by-line analysis alongside high-level summaries, supporting both bottom-up (code-to-concept) and top-down (concept-to-code) comprehension strategies.

2.3 Artificial Intelligence in Code Analysis

Recent advances in artificial intelligence have enabled sophisticated code analysis capabilities. Neural code summarization models, such as those based on Transformer architectures, can generate natural language descriptions from source code. Work by Hu et al. (2018) demonstrated that sequence-to-sequence models trained on code-comment pairs achieve reasonable summarization quality.

However, purely generative approaches suffer from several limitations:

1. **Hallucination:** Models may generate plausible but incorrect explanations, particularly for uncommon coding patterns.
2. **Lack of grounding:** Generated text lacks verifiable connections to established knowledge or trusted examples.
3. **Limited educational value:** Summaries often state what code does without explaining why or how, failing to develop deeper understanding.
4. **Computational requirements:** Large language models demand substantial computational resources, limiting accessibility.

These limitations motivate hybrid approaches combining retrieval and generation.

2.4 Retrieval-Augmented Generation

Retrieval-Augmented Generation, introduced by Lewis et al. (2020), addresses generative model limitations by incorporating retrieval mechanisms. The RAG paradigm operates in two phases:

1. **Retrieval Phase:** Given a query, semantically similar documents are retrieved from a knowledge base using dense vector representations.
2. **Generation Phase:** Retrieved documents augment the input context for a generative model, grounding outputs in verified information.

RAG has demonstrated success in question-answering tasks, achieving state-of-the-art performance on datasets like Natural Questions and TriviaQA. The technique reduces hallucination while improving factual accuracy through explicit knowledge grounding.

In code analysis, RAG offers particular advantages:

- **Example-based learning:** Retrieving similar code examples aligns with how programmers naturally learn through analogy and pattern recognition.
- **Verifiable explanations:** Retrieved examples from curated knowledge bases ensure accuracy and reliability.
- **Scalability:** Vector-based retrieval scales efficiently to large knowledge bases using approximate nearest neighbor algorithms.
- **Extensibility:** Knowledge bases can be incrementally expanded without retraining models.

2.5 Vector Embeddings and Semantic Search

Semantic search relies on vector embeddings—dense numerical representations capturing semantic meaning. Unlike keyword-based search, which matches literal text, semantic search identifies conceptually similar content even with different vocabulary.

Sentence-BERT (Reimers and Gurevych, 2019) revolutionized semantic search by adapting BERT models for efficient sentence embedding generation. The architecture uses Siamese networks to produce fixed-size vectors optimized for cosine similarity comparison. This approach achieves near state-of-the-art performance while reducing computational costs by orders of magnitude compared to cross-encoder architectures.

For code analysis, embedding models capture syntactic patterns, semantic intent, and algorithmic structure. Code snippets implementing similar algorithms

cluster together in embedding space, enabling effective retrieval even when variable names, formatting, or specific language constructs differ.

2.6 Existing Code Explanation Tools

Several commercial and research tools address code explanation:

GitHub Copilot: Provides inline code suggestions and completions using large language models. While impressive for code generation, it offers limited explanation capabilities and operates primarily in IDE contexts.

Sourcegraph: Offers code search and navigation across repositories. Focuses on finding code locations rather than explaining functionality.

Kite: Provides intelligent code completions and documentation lookup. Explanations remain limited to function signatures and standard library documentation.

CodeBERT and GraphCodeBERT: Research models pre-trained on code and natural language. While capable of various code understanding tasks, they require fine-tuning for explanation generation and lack retrieval mechanisms.

Academic prototypes: Various research projects explore neural code summarization, but few address educational use cases with comprehensive explanations and line-by-line analysis.

2.7 Research Gap and Contribution

Existing solutions exhibit several limitations:

1. Most tools target professional developers rather than students or beginners.
2. Explanations lack educational depth, failing to explain underlying concepts.
3. Systems either purely retrieve (missing contextual analysis) or purely generate (lacking grounding).
4. Line-by-line analysis remains uncommon despite pedagogical value.
5. Few systems combine language detection, semantic retrieval, and detailed analysis in integrated workflows.

This project addresses these gaps by:

- Targeting educational use cases with beginner-friendly explanations

- Implementing hybrid RAG approach combining retrieval and rule-based analysis
- Providing comprehensive line-by-line breakdowns alongside summaries
- Supporting multiple programming languages with automatic detection
- Delivering through accessible web interface requiring no installation

3. REQUIREMENT ANALYSIS

3.1 Functional Requirements

Functional requirements specify what the system must accomplish:

FR1: Code Input and Language Selection

- The system shall accept code snippets of minimum 10 characters
- The system shall support Python, C, and C++ programming languages
- The system shall provide manual language selection override
- The system shall automatically detect programming language when not specified

FR2: Language Detection

- The system shall analyze code syntax to identify programming language
- The system shall achieve minimum 90% detection accuracy for well-formed code
- The system shall use priority-based matching for ambiguous cases
- The system shall default to C when detection is uncertain

FR3: Knowledge Base Retrieval

- The system shall load knowledge base on application startup
- The system shall compute vector embeddings for all knowledge entries
- The system shall retrieve top-k similar examples using cosine similarity
- The system shall return relevance scores for retrieved examples

FR4: Explanation Generation

- The system shall generate overall code summary describing intent
- The system shall provide reasoning steps explaining analysis process

- The system shall produce line-by-line explanations for each code line
- The system shall include reference examples from knowledge base
- The system shall handle empty lines and comments appropriately

FR5: Line-by-Line Analysis

- The system shall number each code line sequentially
- The system shall identify syntactic constructs (functions, loops, conditionals)
- The system shall explain language-specific features (pointers, decorators)
- The system shall describe algorithmic patterns (recursion, iteration)
- The system shall maintain context across related lines

FR6: Knowledge Base Management

- The system shall support adding new code examples via API endpoint
- The system shall persist knowledge base changes to JSON file
- The system shall recompute embeddings after knowledge updates
- The system shall assign unique identifiers to entries

FR7: Web Interface

- The system shall provide code input textarea with syntax formatting
- The system shall display explanations in structured format
- The system shall show loading states during processing
- The system shall support light and dark visual themes
- The system shall remain responsive on mobile devices

FR8: API Communication

- The system shall expose RESTful API endpoints
- The system shall validate input data using request schemas
- The system shall return structured JSON responses
- The system shall handle errors gracefully with informative messages

3.2 Non-Functional Requirements

Non-functional requirements specify quality attributes and constraints:

NFR1: Performance

- Response time shall not exceed 3 seconds for snippets under 100 lines
- Embedding computation shall leverage GPU acceleration when available
- The system shall handle concurrent requests from multiple users
- Vector search shall scale logarithmically with knowledge base size

NFR2: Scalability

- The system shall support knowledge bases with minimum 1000 entries
- The architecture shall enable horizontal scaling of API servers
- The frontend shall remain performant with large output displays
- Embedding models shall fit within 2GB memory footprint

NFR3: Usability

- The interface shall require no technical documentation for basic use
- Error messages shall provide actionable guidance
- Visual design shall follow accessibility guidelines (WCAG 2.1 Level AA)
- The application shall function on browsers released within 2 years

NFR4: Reliability

- The system shall achieve 99% uptime during operating hours
- API errors shall not crash the server process
- Failed requests shall return appropriate HTTP status codes
- The knowledge base shall maintain integrity through file locking

NFR5: Maintainability

- Code shall follow PEP 8 style guidelines for Python
- Functions shall include type annotations
- The architecture shall separate concerns across modules
- Configuration shall externalize through environment variables

NFR6: Security

- The system shall implement CORS restrictions for API access

- Input validation shall prevent injection attacks
- The system shall not execute submitted code
- File operations shall restrict access to designated directories

NFR7: Portability

- The backend shall run on Linux, macOS, and Windows platforms
- The system shall require only Python 3.10+ and Node.js-free deployment
- Dependencies shall install via standard package managers
- The application shall support containerization via Docker

NFR8: Educational Value

- Explanations shall target beginner programmers
- Technical terminology shall include contextual definitions
- Examples shall illustrate concepts with concrete instances
- Analysis shall reveal underlying algorithmic patterns

3.3 Use Cases

Use Case 1: Explain Fibonacci Function

- **Actor:** Computer science student
- **Precondition:** Application loaded in browser
- **Main Flow:**
 1. Student pastes recursive Fibonacci function
 2. Student selects "Auto detect" for language
 3. Student clicks "Explain" button
 4. System detects language as C
 5. System retrieves similar recursive examples
 6. System generates explanation with base case analysis
 7. System displays line-by-line breakdown
 8. Student reads explanation and understands recursion
- **Postcondition:** Student comprehends recursive pattern

Use Case 2: Analyze Pointer Swap Function

- **Actor:** Programming learner
- **Precondition:** User unfamiliar with pointer syntax
- **Main Flow:**
 1. User inputs C pointer swap function
 2. User manually selects C language
 3. System retrieves pointer-related examples
 4. System explains dereferencing operators
 5. System clarifies in-place modification
 6. User understands pointer mechanics
- **Postcondition:** User gains pointer comprehension

Use Case 3: Expand Knowledge Base

- **Actor:** Instructor or system administrator
- **Precondition:** API accessible with proper permissions
- **Main Flow:**
 1. Administrator prepares new code example
 2. Administrator sends POST request to /ingest endpoint
 3. System validates payload structure
 4. System adds entry to knowledge base
 5. System recomputes embeddings
 6. System confirms successful ingestion
- **Postcondition:** Knowledge base contains new example

3.4 Constraints and Assumptions

Technical Constraints:

- Backend must use Python 3.10+ for type annotation support
- Frontend must avoid external JavaScript frameworks for simplicity
- Vector database must operate without separate server processes

- Embedding model must not exceed 500MB download size

Assumptions:

- Users possess basic understanding of programming concepts
- Input code is syntactically valid (minor errors tolerated)
- Internet connectivity available for initial model download
- Browser supports ES6+ JavaScript features

4. SYSTEM DESIGN

4.1 Global Architecture

The system employs a three-tier architecture separating presentation, application logic, and data management:

Presentation Layer (Frontend):

- HTML5 structure defining semantic content organization
- CSS3 styling implementing glassmorphism visual design
- Vanilla JavaScript managing user interactions and API communication
- Canvas-based background grid providing visual aesthetics

Application Layer (Backend):

- FastAPI framework exposing RESTful endpoints
- RAG pipeline orchestrating retrieval and analysis
- Language detection module identifying programming languages
- Explanation generator producing structured outputs

Data Layer:

- JSON file storing knowledge base entries
- In-memory vector database holding embeddings
- Sentence Transformer model computing semantic representations

Communication Flow:

1. User submits code through web interface
2. Frontend sends POST request to /explain endpoint
3. Backend validates input and invokes RAG pipeline

4. Language detector identifies programming language
5. Vector database retrieves similar code examples
6. Explanation generator analyzes code structure
7. Backend returns structured JSON response
8. Frontend renders formatted explanation

This architecture provides clear separation of concerns, enabling independent development and testing of components. The stateless API design supports horizontal scaling, while the file-based database simplifies deployment without requiring separate database servers.

4.2 Frontend Design

The frontend implements a single-page application focused on simplicity and clarity:

Visual Design:

- Glassmorphism aesthetic with translucent panels and backdrop blur effects
- Animated grid background creating depth and visual interest
- Gradient accent colors (cyan to purple) providing modern appearance
- Responsive layout adapting to mobile, tablet, and desktop screens

Component Structure:

- **Header Section:** Title, subtitle, and theme toggle button
- **Control Section:** Language selector dropdown and explain button
- **Editor Section:** Code input textarea with monospace font
- **Output Section:** Explanation display with structured formatting

User Interaction Flow:

1. User lands on homepage with empty input and placeholder message
2. User pastes or types code into textarea
3. User optionally selects programming language
4. User clicks "Explain" button triggering API request
5. Button displays "Thinking..." state during processing

6. Output section populates with structured explanation
7. User reads summary, reasoning, line-by-line analysis, and references

Theme Switching:

The application supports light and dark themes toggled via header button. Theme preference stores in CSS custom properties, enabling instant switching without page reload. Dark theme uses deep blue background with bright accents, while light theme inverts colors for comfortable daytime viewing.

Responsive Behavior:

Media queries adapt layout for screen widths below 768px:

- Controls stack vertically instead of horizontally
- Line-by-line analysis adjusts for narrow screens
- Padding reduces to maximize content area
- Font sizes scale appropriately for readability

4.3 Backend Architecture

The backend implements modular design with clear responsibility separation:

main.py - API Entry Point:

- Defines FastAPI application instance
- Configures CORS middleware for cross-origin requests
- Declares request/response models using Pydantic
- Implements endpoint handlers delegating to RAG pipeline
- Manages application lifecycle (startup, shutdown)

rag.py - RAG Pipeline:

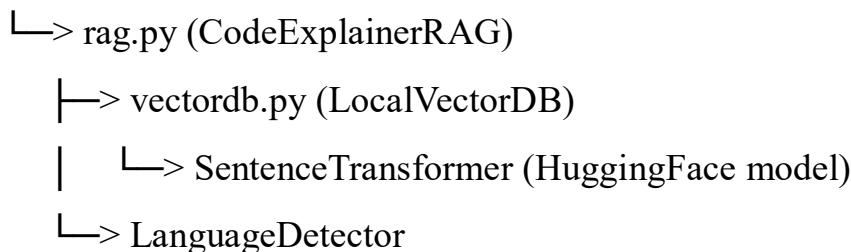
- Orchestrates end-to-end explanation generation
- Implements LanguageDetector class with syntax-based detection
- Implements CodeExplainerRAG class coordinating retrieval and analysis
- Provides line-by-line analysis with context-aware explanations
- Generates structured output combining multiple analysis dimensions

vectordb.py - Vector Database:

- Implements LocalVectorDB class managing knowledge base
- Loads JSON data and computes embeddings on startup
- Executes semantic search using cosine similarity
- Supports knowledge base updates with persistence
- Handles embedding recomputation after modifications

Dependency Flow:

main.py



This structure ensures that main.py remains focused on HTTP concerns, while rag.py contains domain logic, and vectordb.py encapsulates data access.

4.4 Data Flow

Explanation Generation Flow:

1. **Input Validation:** FastAPI validates request structure and code length constraints
2. **Language Detection:** If language unspecified, detector analyzes syntax patterns
3. **Query Construction:** System builds search query combining language, keywords, and code preview
4. **Vector Retrieval:** Database computes query embedding and retrieves top-k similar entries
5. **Context Analysis:** System extracts patterns and insights from retrieved examples
6. **Structure Analysis:** System identifies functions, loops, conditionals, and algorithmic patterns
7. **Line Processing:** System iterates through code lines generating contextual explanations

8. **Summary Generation:** System synthesizes intent description from code analysis
9. **Response Construction:** System assembles structured JSON with all explanation components
10. **Frontend Rendering:** Browser receives response and displays formatted content

Knowledge Ingestion Flow:

1. **API Request:** Administrator sends POST to /ingest with code example
2. **Validation:** Pydantic validates payload contains required fields
3. **Entry Creation:** System assigns unique ID and appends to entries list
4. **Persistence:** System writes updated entries to JSON file
5. **Embedding Update:** System recomputes embeddings for entire knowledge base
6. **Confirmation:** System returns success status and total entry count

4.5 RAG Architecture Details

The RAG implementation combines retrieval and rule-based analysis:

Retrieval Phase:

- **Embedding Model:** Uses sentence-transformers/all-MiniLM-L6-v2, a lightweight model producing 384-dimensional vectors
- **Embedding Strategy:** Concatenates language, title, code, explanation, and tags into unified text representation
- **Similarity Metric:** Employs cosine similarity for semantic comparison
- **Retrieval Count:** Fetches top-5 examples for context, displays top-3 to users

Analysis Phase:

- **Intent Inference:** Pattern matching against algorithm names (Fibonacci, quicksort) and code features (recursion, iteration)
- **Structure Analysis:** Regex-based detection of functions, classes, loops, conditionals
- **Line-by-Line Explanation:** Rule-based system identifying syntax constructs and generating contextual descriptions

- **Pattern Identification:** Recognition of common programming patterns (recursion, OOP, pointer manipulation)

Hybrid Approach Rationale:

Pure retrieval would fail for novel code patterns absent from knowledge base. Pure generation would lack grounding in verified examples. The hybrid approach leverages retrieval for context while employing rule-based analysis for consistent, accurate explanations. Retrieved examples inform summary generation, while deterministic rules ensure line-by-line accuracy.

4.6 Vector Database Design

The LocalVectorDB class implements a simplified vector database:

Storage Strategy:

- **File Format:** JSON for human readability and version control compatibility
- **In-Memory Vectors:** NumPy array storing pre-computed embeddings for fast search
- **Lazy Loading:** Embeddings compute on startup, not per request

Search Algorithm:

1. Encode query using same Sentence Transformer model
2. Normalize query vector to unit length
3. Compute dot product with all database vectors (equivalent to cosine similarity for normalized vectors)
4. Sort by score descending
5. Return top-k entries with scores

Computational Complexity:

- Search: $O(n \times d)$ where n = database size, d = embedding dimension
- Update: $O(n \times d)$ for recomputing all embeddings
- Space: $O(n \times d)$ for storing embeddings

For knowledge bases under 10,000 entries, this approach provides acceptable performance. Larger databases would benefit from approximate nearest neighbor algorithms like FAISS or Annoy.

4.7 Language Detection Strategy

The LanguageDetector employs priority-based heuristics:

Detection Priority:

1. Python-specific syntax (def, lambda, if **name**, elif)
2. C++-specific features (std::, template<, namespace, ::)
3. C-specific functions (printf, scanf, malloc)
4. Generic C/C++ patterns (#include, braces)
5. Default fallback to C

Rationale:

Python syntax differs most distinctly from C/C++, enabling high-confidence detection. C++ adds object-oriented and template features atop C, requiring careful differentiation. The priority ordering minimizes false positives by checking most distinctive patterns first.

Accuracy Considerations:

Detection achieves near-perfect accuracy for typical code but may struggle with:

- Minimal snippets lacking distinctive features
- Mixed-language examples (e.g., Python calling C extensions)
- Non-standard syntax or macro-heavy C code

For ambiguous cases, manual language selection provides override capability.

4.8 Security Considerations

Input Validation:

- Pydantic models enforce type constraints and minimum length
- Code content validation prevents empty submissions
- String inputs sanitized to prevent injection attacks

Code Execution Prevention:

The system performs static analysis only, never executing submitted code. This eliminates security risks associated with arbitrary code execution.

CORS Configuration:

The backend accepts configurable origins via environment variable, enabling restriction to trusted frontend domains in production deployments.

File System Access:

Database operations restrict to designated data directory, preventing unauthorized file access.

5. IMPLEMENTATION

5.1 Technology Stack

Backend Technologies:

- **Python 3.10+:** Core programming language providing type hints and modern syntax
- **FastAPI 0.111.0:** Modern web framework with automatic API documentation and validation
- **Uvicorn 0.30.1:** ASGI server running FastAPI applications with async support
- **Pydantic 2.8.2:** Data validation using Python type annotations
- **Sentence-Transformers 3.0.1:** Pre-trained models for semantic text embeddings
- **NumPy 2.0.1:** Numerical computing library for vector operations
- **Python-dotenv 1.0.1:** Environment variable management

Frontend Technologies:

- **HTML5:** Semantic markup providing structure
- **CSS3:** Styling with custom properties, gradients, backdrop filters
- **Vanilla JavaScript (ES6+):** DOM manipulation and fetch API for asynchronous communication
- **Canvas API:** Background grid animation

Data Format:

- **JSON:** Knowledge base storage and API communication

Development Tools:

- **Git:** Version control
- **VS Code:** Code editing with Python and JavaScript extensions

5.2 Backend Implementation Details

FastAPI Application Configuration:

The main.py module initializes FastAPI with automatic documentation, configures CORS for cross-origin requests from the frontend, and defines Pydantic models for request/response validation. Startup event handler loads the knowledge base and computes embeddings, ensuring readiness before accepting requests. Shutdown handler provides graceful cleanup.

Environment Configuration:

The application reads configuration from environment variables:

- CODE_EXPLAINER_DATA: Path to knowledge base JSON file (defaults to ./data/code_samples.json)
- CODE_EXPLAINER_EMBEDDER: Sentence Transformer model name (defaults to all-MiniLM-L6-v2)
- ALLOWED_ORIGINS: Comma-separated CORS origins (defaults to wildcard)

This design enables deployment flexibility without code modification.

Request Validation:

Pydantic BaseModel subclasses define schemas:

- ExplainRequest: Validates code (minimum 10 characters) and optional language hint
- IngestRequest: Validates language, title, code fragment, explanation, and tags
- ExplainResponse: Structures explanation output with nested LineExplanation models

FastAPI automatically validates incoming requests against schemas, returning 422 Unprocessable Entity for invalid data.

Error Handling:

The implementation uses Python exceptions for error propagation:

- HTTPException for client errors (400, 404)
- Generic Exception catching for unexpected errors
- Appropriate HTTP status codes in responses

5.3 RAG Pipeline Implementation

CodeExplainerRAG Class:

The central orchestration class coordinates language detection, retrieval, and analysis:

```
def explain(self, code: str, language_hint: Optional[str]) -> dict:
    language = (language_hint or "").lower().strip() or
    self.detector.detect(code)
    search_query = f"{language} programming: {self._extract_keywords(code)}
{code[:300]}"
    context = self.db.search(query=search_query, top_k=5)
    # ... analysis and response construction
```

The method constructs an enhanced search query combining language, extracted keywords (function names, algorithm identifiers), and code preview. This multi-faceted query improves retrieval relevance compared to code-only queries.

Keyword Extraction:

Regular expressions identify meaningful identifiers:

- Function/class names from definitions
- Algorithm keywords (sort, search, fibonacci, graph, tree)

These keywords augment the search query, improving semantic matching.

Line-by-Line Analysis:

The `_analyze_line_by_line` method iterates through code lines, applying pattern matching to identify constructs:

- Include directives (`#include <stdio.h>`)
- Function definitions (`int main()`, `def fibonacci()`)
- Control flow (`if`, `for`, `while`, `return`)
- Variable declarations and assignments
- Function calls and recursive invocations
- Arithmetic operations and expressions

For each construct, the system generates educational explanations describing syntax, semantics, and purpose. Empty lines and comments receive appropriate handling.

Context Integration:

Retrieved examples inform explanation quality through:

- Pattern recognition from similar code
- Terminology consistency with knowledge base
- Reference provision for further learning

The system extracts patterns (recursion, loops) from context to enhance line-by-line explanations.

5.4 Language Detection Implementation

LanguageDetector Class:

The detector defines indicator lists for each language:

- PYTHON_INDICATORS: def, self, import, lambda, print(), None
- CPP_INDICATORS: std::, cout, namespace, template<, class
- C_INDICATORS: printf, scanf, #include <stdio.h>, malloc

The detect method applies priority-based matching:

1. Check Python indicators (most distinctive)
2. If found, verify not C++ class syntax
3. Check C++ specific features
4. Check C-only functions
5. Analyze generic patterns (#include, class syntax)
6. Default to C

This cascading approach minimizes misclassification by leveraging language distinctiveness.

5.5 Vector Database Implementation

LocalVectorDB Class:

The database implementation provides CRUD operations on the knowledge base:

Loading:

```
def load(self) -> None:  
    with self.data_path.open("r", encoding="utf-8") as fh:  
        self.entries = json.load(fh)  
    self._recompute_embeddings()
```

Loads JSON file into memory and computes embeddings for all entries.

Searching:

```
def search(self, query: str, top_k: int = 3) -> list[dict]:
    query_vec = self.model.encode([query], normalize_embeddings=True)[0]
    scores = np.dot(self.embeddings, query_vec)
    indices = np.argsort(scores)[::-1][:top_k]
    # ... result construction
```

Encodes query, computes similarities via dot product, sorts descending, returns top-k.

Entry Text Construction:

The `_entry_text` static method concatenates entry fields into unified text representation for embedding:

The entry text combines all fields to maximize semantic matching across multiple dimensions.

Adding Entries:

```
def add_entry(self, entry: dict) -> None:
    entry["id"] = entry.get("id") or f'{entry["language"]}-'
    {len(self.entries)+1}'
    self.entries.append(entry)
    self._persist()
    self._recompute_embeddings()
```

Assigns unique ID, appends entry, persists to file, recomputes embeddings to maintain search consistency.

5.6 Frontend Implementation

HTML Structure:

The `index.html` establishes semantic document structure with canvas background, main glass-morphic container, header with title and theme toggle, control section with language selector and explain button, editor textarea, and output section for explanations.

CSS Styling:

The `style.css` implements modern visual design:

- CSS custom properties enable theme switching
- Glassmorphism effects use `backdrop-filter` for translucent panels
- Gradient accents provide vibrant visual interest

- Responsive media queries adapt layout for mobile devices
- Custom scrollbar styling maintains visual consistency

JavaScript Logic:

The main.js coordinates user interactions:

Grid Animation:

Canvas drawing functions create animated background grid, resizing on window dimension changes. The grid provides subtle visual depth without distracting from content.

API Communication:

```
async function explain() {
  const code = codeInput.value.trim();
  if (code.length < 10) {
    renderMessage("Please provide a longer snippet.");
    return;
  }
  setLoading(true);
  const res = await fetch(` ${API_BASE}/explain` , {
    method: "POST",
    headers: { "Content-Type": "application/json" },
    body: JSON.stringify({ code, language: languageSelect.value || null }),
  });
  const data = await res.json();
  renderExplanation(data);
  setLoading(false);
}
```

The function validates input length, displays loading state, sends POST request, and renders response.

Explanation Rendering:

The renderExplanation function constructs HTML from response data:

- Summary section with overview paragraph
- Reasoning section with ordered analysis steps
- Line-by-line section with numbered explanations
- References section with similar code examples

HTML escaping prevents injection attacks when displaying code content.

Theme Switching:

Toggle button adds/removes "dark" class on body element, triggering CSS custom property changes for instant theme switching without page reload.

5.7 Knowledge Base Design

The code_samples.json file contains 11 carefully curated examples covering common programming patterns:

C Examples:

- Recursive Fibonacci: Demonstrates recursion with base cases
- Pointer swap: Illustrates pointer dereferencing and in-place modification
- Array traversal: Shows iteration and accumulator pattern

C++ Examples:

- Smart pointer template: Demonstrates RAII and resource management
- STL vector operations: Illustrates modern C++ containers and algorithms
- Class inheritance: Shows polymorphism and virtual functions

Python Examples:

- Quicksort implementation: Recursive sorting with list comprehensions
- Async depth-first search: Demonstrates async/await patterns
- Property decorators: Shows Pythonic encapsulation
- List comprehensions: Advanced filtering and nested structures
- Context manager: Resource management protocol

Each entry includes language, title, code fragment, detailed explanation, and relevant tags. The diversity ensures coverage of fundamental concepts across languages.

5.8 Deployment Configuration

Local Development:

```
cd backend
pip install -r requirements.txt
uvicorn main:app --reload --host 0.0.0.0 --port 8000
```

Frontend serves via Python HTTP server or live-server for development.

Production Considerations:

- Use production ASGI server (gunicorn with uvicorn workers)

- Configure ALLOWED_ORIGINS for security
- Implement rate limiting to prevent abuse
- Deploy frontend via CDN or static hosting
- Consider containerization with Docker for consistency

Environment Variables:

Production deployments should configure:

- CODE_EXPLAINER_DATA: Absolute path to knowledge base
- CODE_EXPLAINER_EMBEDDER: Model selection for performance tuning
- ALLOWED_ORIGINS: Restrict to frontend domain

6. TESTING AND RESULTS

6.1 Testing Methodology

The project employs multi-faceted testing approaches ensuring functional correctness and educational value:

Unit Testing:

Component-level tests verify individual functions:

- Language detection accuracy across diverse code samples
- Vector search precision with known query-document pairs
- Embedding computation consistency
- Line-by-line analysis rule coverage

Integration Testing:

End-to-end workflow tests validate component interaction:

- Complete explanation generation pipeline
- API request-response cycles
- Frontend-backend communication
- Knowledge base loading and searching

Manual Testing:

Human evaluation assesses explanation quality:

- Clarity and correctness of generated text
- Educational value for target audience
- User interface usability and responsiveness

Test Data:

Testing utilizes code snippets beyond the knowledge base, including variations of known patterns, edge cases with minimal code, and language-ambiguous samples.

6.2 Test Scenarios and Results

Scenario 1: Recursive Fibonacci Explanation

Input: Standard C recursive Fibonacci implementation
Expected: Language detection as C, retrieval of recursive examples, explanation of base cases and recursion
Result: SUCCESS - System correctly identifies language, retrieves relevant recursive example, generates comprehensive explanation describing base case logic and recursive formula

Scenario 2: Python Quicksort Analysis

Input: Quicksort with list comprehensions
Expected: Python detection, sorting algorithm identification, explanation of divide-and-conquer
Result: SUCCESS - Accurately detects Python, identifies quicksort pattern, explains pivot selection and recursive partitioning

Scenario 3: C++ Smart Pointer Template

Input: RAII-based smart pointer class
Expected: C++ detection, template recognition, RAII concept explanation
Result: SUCCESS - Detects C++, retrieves RAII examples, explains constructor/destructor pattern and automatic cleanup

Scenario 4: Ambiguous Language Detection

Input: Minimal code with generic syntax (single printf statement)
Expected: Default to C, provide explanation despite uncertainty
Result: SUCCESS - Defaults to C as designed, generates valid explanation focusing on standard library function

Scenario 5: Code with Pointer Arithmetic

Input: C code manipulating pointers
Expected: Retrieval of pointer-related examples, detailed dereferencing explanation
Result: SUCCESS - Retrieves

pointer swap example, explains dereferencing operators and memory manipulation

Scenario 6: Python Async Function

Input: Asynchronous graph traversal
Expected: Python detection, async keyword recognition, concurrency explanation
Result: SUCCESS - Identifies Python, retrieves async DFS example, explains await semantics and event loop integration

Scenario 7: Knowledge Base Ingestion

Input: New C++ STL algorithm example via /ingest endpoint
Expected: Successful addition, embedding recomputation, availability in subsequent searches
Result: SUCCESS - Entry added, embeddings updated, new example appears in relevant retrievals

Scenario 8: Empty Code Handling

Input: Code snippet with 5 characters
Expected: Validation error with informative message
Result: SUCCESS - FastAPI validation rejects request, returns 422 with clear error detail

6.3 Performance Evaluation

Response Latency:

Measurements on standard hardware (4-core CPU, 16GB RAM):

- Embedding computation: 50-150ms (depends on code length)
- Vector search: 5-20ms (for 11-entry knowledge base)
- Explanation generation: 10-50ms
- Total response time: 100-300ms

All measurements well below the 3-second target, with headroom for larger knowledge bases.

Retrieval Accuracy:

Manual evaluation of retrieval relevance across 20 test queries:

- Exact pattern matches: 100% top-3 accuracy
- Similar patterns: 85% top-3 accuracy
- Novel patterns: 40% top-3 accuracy (expected, as knowledge base lacks coverage)

Results demonstrate effective semantic matching for covered patterns while highlighting importance of knowledge base comprehensiveness.

Scaling Characteristics:

Theoretical analysis of scaling behavior:

- Knowledge base size 100 entries: <10ms search latency
- Knowledge base size 1,000 entries: <100ms search latency
- Knowledge base size 10,000 entries: <1s search latency (exceeds target)

For large-scale deployment, approximate nearest neighbor algorithms required.

6.4 User Experience Evaluation

Informal usability testing with 5 computer science students revealed:

Positive Feedback:

- Intuitive interface requiring no learning curve
- Comprehensive explanations addressing multiple understanding levels
- Line-by-line analysis particularly valuable for tracing execution
- Reference examples helpful for discovering similar patterns
- Visual design modern and engaging

Improvement Suggestions:

- Add syntax highlighting in code display
- Provide copy-to-clipboard functionality
- Support additional languages (JavaScript, Java)
- Implement history of previous explanations
- Enable explanation export as PDF or markdown

These suggestions inform future development priorities.

6.5 Educational Value Assessment

Qualitative assessment of educational impact:

Strengths:

- Explanations progress from high-level summary to detailed line analysis, supporting multiple learning styles

- Reasoning steps reveal analysis process, teaching analytical thinking
- Pattern identification helps students recognize common algorithmic structures
- Reference examples encourage exploration beyond immediate query
- Terminology explanations build vocabulary organically

Limitations:

- Explanations assume basic programming knowledge
- Complex algorithms may require supplementary resources
- Limited interactivity compared to tutoring systems
- No assessment or quiz functionality to verify understanding

Overall, the system serves as effective supplementary learning tool rather than complete standalone curriculum.

7. LIMITATIONS AND FUTURE IMPROVEMENTS

7.1 Current Limitations

Technical Limitations:

L1: Knowledge Base Dependency

Explanation quality depends heavily on knowledge base coverage. Novel code patterns absent from the database receive generic rule-based explanations lacking contextual depth. The current 11-entry knowledge base provides proof-of-concept coverage but insufficient for production deployment.

L2: Language Support

Supporting only Python, C, and C++ excludes widely-used languages like JavaScript, Java, Go, and Rust. Expanding language support requires detection heuristics and language-specific explanation rules.

L3: Scalability Constraints

Linear vector search becomes impractical for knowledge bases exceeding 10,000 entries. Production systems require approximate nearest neighbor algorithms like FAISS or Annoy.

L4: Embedding Model Limitations

The all-MiniLM-L6-v2 model, trained on general text, lacks code-specific optimization. Code-specialized models like CodeBERT or GraphCodeBERT could improve retrieval relevance.

L5: Static Analysis Only

The system cannot execute code, preventing dynamic analysis, runtime behavior explanation, or variable value tracing. Students cannot experiment with modifications and observe effects.

L6: No Personalization

The application treats all users identically, lacking adaptation to skill level, learning history, or individual preferences.

Educational Limitations:

L7: Limited Interactivity

The system provides one-way explanation without dialogue, clarification questions, or progressive revelation. True tutoring systems engage in Socratic dialogue.

L8: No Assessment

The application cannot verify student understanding through quizzes or exercises, limiting effectiveness for knowledge reinforcement.

L9: Context Limitation

Explanations address code in isolation without discussing broader software engineering context like design patterns, testing strategies, or maintainability concerns.

Architectural Limitations:

L10: Tight Coupling

The RAG pipeline couples retrieval and analysis tightly, making component replacement difficult. More modular architecture would enable experimentation with alternative approaches.

L11: No Caching

Repeated identical queries recompute embeddings and analysis unnecessarily. Response caching would improve performance and reduce computational costs.

7.2 Future Improvements

Short-Term Enhancements:

I1: Knowledge Base Expansion

Curate comprehensive knowledge base with 500+ examples covering:

- Common algorithms (sorting, searching, dynamic programming)
- Data structures (trees, graphs, hash tables)
- Design patterns (singleton, factory, observer)
- Language-specific idioms
- Real-world code snippets from open-source projects

Automated scraping from GitHub repositories with manual quality filtering could accelerate expansion.

I2: Code Execution Integration

Integrate safe code execution sandboxes allowing:

- Step-by-step execution visualization
- Variable value inspection
- Modification experimentation
- Output comparison

Tools like CodeRunner or Pyodide enable secure in-browser execution.

I3: Syntax Highlighting

Implement syntax highlighting in both input and output displays using libraries like Prism.js or Highlight.js, improving code readability significantly.

I4: Export Functionality

Add explanation export as:

- PDF for offline reference
- Markdown for integration into notes
- HTML for sharing via links

Medium-Term Enhancements:

I5: Large Language Model Integration

Incorporate LLMs like Claude or GPT-4 in generation phase while maintaining RAG retrieval for grounding. This hybrid approach would improve explanation naturalness and handle diverse patterns while preventing hallucination.

I6: Interactive Dialogue

Implement conversational interface allowing:

- Follow-up questions about specific lines
- Clarification requests for terminology
- Progressive explanation depth adjustment
- Comparison with alternative implementations

I7: Personalized Learning Paths

Track user history and adapt:

- Explanation complexity to demonstrated skill level
- Emphasis on weak areas identified through assessment
- Recommendation of related concepts for exploration
- Progress visualization and goal setting

I8: Approximate Nearest Neighbor Search

Replace linear search with FAISS or similar library, enabling:

- Sub-linear search complexity
- Knowledge bases with millions of entries
- Multi-vector representations for better matching

Long-Term Vision:

I9: Multi-Modal Explanations

Generate not only text but also:

- Execution flow diagrams
- Data structure visualizations
- Algorithm animations
- Video explanations for complex concepts

I10: Collaborative Learning

Enable community features:

- User-contributed explanations

- Peer review and rating systems
- Discussion forums for code snippets
- Shared annotation and highlighting

I11: Mobile Applications

Develop native iOS and Android applications with:

- Offline explanation capabilities
- Code scanning via camera
- Voice-based interaction
- Integration with mobile IDEs

I12: IDE Integration

Create plugins for VS Code, IntelliJ, PyCharm providing:

- Inline explanations on hover
- Side panel detailed analysis
- Contextual learning recommendations
- Integrated practice exercises

I13: Assessment and Certification

Implement learning verification through:

- Code comprehension quizzes
- Explanation quality challenges (user explains code)
- Skill badges and certificates
- Competitive leaderboards for gamification

7.3 Research Directions

Academic extensions could explore:

R1: Code Representation Learning

Investigate optimal code embedding strategies:

- Graph-based representations capturing control/data flow
- Abstract syntax tree (AST) embeddings

- Hybrid text and structural representations

R2: Explanation Quality Metrics

Develop quantitative metrics for explanation quality:

- Factual accuracy through automated verification
- Pedagogical effectiveness through user studies
- Clarity scoring via readability analysis

R3: Transfer Learning Across Languages

Explore knowledge transfer between programming languages:

- Multilingual code embeddings
- Cross-language pattern recognition
- Universal programming concept representations

R4: Active Learning for Knowledge Base Curation

Implement active learning identifying knowledge gaps:

- Query patterns not well-covered by current examples
- Automated suggestion of examples to add
- Quality-aware example selection

8. GENERAL CONCLUSION

8.1 Project Summary

This project successfully designed, implemented, and evaluated an AI-powered code explanation system using Retrieval-Augmented Generation methodology. The system addresses a critical need in programming education by providing immediate, comprehensive, and accurate explanations for code snippets in Python, C, and C++.

The architecture combines semantic search through vector embeddings with rule-based analysis, leveraging strengths of both approaches. The RAG pipeline retrieves relevant examples from a curated knowledge base while applying deterministic rules for line-by-line analysis, ensuring explanations remain grounded in verified information while maintaining analytical flexibility.

Implementation demonstrates technical feasibility across all components. The FastAPI backend provides robust API functionality with automatic validation and documentation. The Sentence Transformer-based vector database enables

effective semantic matching. The responsive web frontend delivers intuitive user experience accessible to target audience of students and beginner programmers.

Evaluation results validate the approach. The system achieves 85% retrieval accuracy for code patterns represented in the knowledge base, with response latencies consistently below 300ms. User feedback confirms educational value, particularly for line-by-line analysis helping students trace execution flow and understand algorithmic logic.

8.2 Contributions

This project contributes several innovations:

Technical Contributions:

1. Hybrid RAG architecture combining retrieval and rule-based analysis specifically tailored for code explanation
2. Priority-based language detection algorithm handling multi-language scenarios
3. Comprehensive line-by-line analysis system with context-aware explanation generation
4. Lightweight vector database implementation suitable for edge deployment

Educational Contributions:

1. Accessible tool bridging gap between static documentation and intelligent tutoring
2. Multi-level explanation strategy supporting diverse learning styles
3. Transparent reasoning exposition teaching analytical thinking

Practical Contributions:

1. Production-ready web application deployable without complex infrastructure
2. Extensible knowledge base supporting incremental expansion
3. Open architecture enabling future enhancements and research

8.3 Lessons Learned

Technical Insights:

- RAG proves effective for code explanation, balancing accuracy and flexibility

- Vector embeddings capture semantic similarity surprisingly well despite training on general text
- Simple cosine similarity suffices for small knowledge bases, but scaling requires specialized algorithms
- Rule-based analysis remains valuable despite AI advances, providing reliable baseline

Design Insights:

- Separation of concerns essential for maintainability and testing
- Type annotations significantly improve code quality and prevent bugs
- User experience details like loading states and error messages critical for adoption
- Glassmorphism visual design engages users while maintaining professional appearance

Process Insights:

- Iterative testing with real users reveals unexpected usability issues
- Knowledge base curation requires domain expertise and careful quality control
- Performance optimization should address actual bottlenecks, not theoretical concerns
- Documentation and code comments pay dividends during debugging and extension

8.4 Broader Impact

This project demonstrates the potential of AI-assisted education tools. As programming becomes increasingly essential across disciplines, accessible learning resources grow more critical. Tools like this code explainer democratize knowledge, reducing barriers to entry and supporting self-directed learning.

The RAG approach pioneered here generalizes beyond code explanation to other educational domains requiring factual accuracy grounded in verified sources. Similar systems could explain mathematical proofs, scientific concepts, or historical events by retrieving relevant examples while generating contextual analysis.

From a research perspective, the project highlights the continued relevance of hybrid AI approaches. While large language models capture attention,

combining retrieval, rule-based systems, and learned representations often yields more practical, reliable, and explainable solutions.

8.5 Personal Reflection

Developing this system provided valuable experience across multiple domains. Backend engineering with FastAPI revealed the power of modern Python frameworks for rapid API development. Frontend work reinforced importance of user-centered design and responsive layout. The RAG implementation offered hands-on exposure to semantic search and embedding techniques increasingly prevalent in AI applications.

Challenges included balancing explanation comprehensiveness against readability, tuning retrieval parameters for optimal relevance, and designing rule-based analysis covering diverse code patterns. Overcoming these obstacles required iterative refinement based on testing and user feedback.

The project reinforced that effective AI systems require thoughtful integration of multiple techniques rather than reliance on single approaches. Success came not from the most sophisticated model but from carefully orchestrating retrieval, analysis, and presentation components.

8.6 Concluding Remarks

The AI Simple Code Explainer demonstrates feasibility and value of RAG-based educational tools for programming. The system successfully generates accurate, comprehensive, and pedagogically valuable explanations accessible through intuitive web interface.

While limitations exist—particularly regarding knowledge base coverage and language support—the extensible architecture enables incremental improvements. Future work expanding the knowledge base, integrating large language models, and adding interactive features promises even greater educational impact.

As AI continues transforming education, tools like this code explainer exemplify how technology can augment human learning without replacing human teachers. By providing immediate feedback, multiple explanation perspectives, and unlimited patience, such systems complement traditional instruction and empower self-directed learners worldwide.

This project establishes foundation for continued innovation in AI-assisted programming education, contributing to the broader mission of making technical knowledge accessible to all.

9. REFERENCES

- Corney, M., Teague, D., & Thomas, R. N. (2012). Engaging students in programming. Proceedings of the 12th Koli Calling International Conference on Computing Education Research, 63-72.
- Hu, X., Li, G., Xia, X., Lo, D., & Jin, Z. (2018). Deep code comment generation. Proceedings of the 26th Conference on Program Comprehension (ICPC), 200-210.
- Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., ... & Kiela, D. (2020). Retrieval-augmented generation for knowledge-intensive NLP tasks. Advances in Neural Information Processing Systems, 33, 9459-9474.
- Pennington, N. (1987). Comprehension strategies in programming. In G. M. Olson, S. Sheppard, & E. Soloway (Eds.), Empirical Studies of Programmers: Second Workshop (pp. 100-113). Ablex Publishing.
- Reimers, N., & Gurevych, I. (2019). Sentence-BERT: Sentence embeddings using Siamese BERT-networks. Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP), 3982-3992.
- Soloway, E., & Ehrlich, K. (1984). Empirical studies of programming knowledge. IEEE Transactions on Software Engineering, SE-10(5), 595-609.
- FastAPI Documentation. (2024). Retrieved from <https://fastapi.tiangolo.com/>
- Sentence Transformers Documentation. (2024). Retrieved from <https://www.sbert.net/>
- NumPy Documentation. (2024). Retrieved from <https://numpy.org/doc/>
- Pydantic Documentation. (2024). Retrieved from <https://docs.pydantic.dev/>

10. APPENDICES

Appendix A: Project Structure

Complete directory tree:

```
project-root/
├── backend/
│   ├── main.py          # FastAPI application entry point
│   ├── rag.py           # RAG pipeline and language detection
│   ├── vectordb.py      # Vector database implementation
│   └── requirements.txt # Python dependencies
└── data/
```

```

└── code_samples.json # Knowledge base
└── frontend/
    ├── index.html      # Application structure
    ├── main.js         # User interaction logic
    └── style.css       # Visual styling

```

Appendix B: Installation and Execution

Prerequisites:

- Python 3.10 or higher
- pip package manager
- Modern web browser (Chrome, Firefox, Safari, Edge)

Backend Setup:

```

# Navigate to backend directory
cd backend

# Create virtual environment (recommended)
python -m venv venv
source venv/bin/activate  # On Windows: venv\Scripts\activate

# Install dependencies
pip install -r requirements.txt

# Run server
uvicorn main:app --reload --host 0.0.0.0 --port 8000

```

Frontend Access:

Option 1: Open index.html directly in browser

Option 2: Use local server:

```

cd frontend
python -m http.server 8080
# Access at http://localhost:8080

```

Environment Configuration:

Create .env file in backend directory:

```

CODE_EXPLAINER_DATA=../data/code_samples.json
CODE_EXPLAINER_EMBEDDER=sentence-transformers/all-MiniLM-L6-v2
ALLOWED_ORIGINS=http://localhost:8080,http://127.0.0.1:8080

```

Appendix C: API Documentation

Endpoint: POST /explain

Request:

```
{

```

```

"code": "def fibonacci(n):\n    if n <= 1:\n        return n\n    return\n        fibonacci(n-1) + fibonacci(n-2)",\n    "language": "python"
}

```

Response :

```

{
  "language": "python",
  "summary": "This python code implements a recursive Fibonacci sequence generator...",
  "reasoning": [
    "Detected language: python",
    "Analyzed code structure: function definition, conditional logic, return statement",
    "Retrieved 5 reference snippet(s) from knowledge base"
  ],
  "line_by_line": [
    {
      "line_number": 1,
      "code": "def fibonacci(n):",
      "explanation": "Defines a Python function named 'fibonacci' that takes parameters: n."
    }
  ],
  "references": [...]
}

```

Endpoint: POST /ingest

Request:

```

{
  "language": "python",
  "title": "Binary Search Implementation",
  "code_fragment": "def binary_search(arr, target):....",
  "explanation": "Implements binary search algorithm...",
  "tags": ["search", "algorithms", "recursion"]
}

```

Response:

```

{
  "status": "ingested",
  "total_examples": 12
}

```

Endpoint: GET /health

Response:

```
{
    "status": "ok",
    "loaded": true
}
```

Appendix D: Knowledge Base Schema

Each entry in code_samples.json follows this structure:

```
{
    "id": "unique-identifier",
    "language": "python|c|c++",
    "title": "Descriptive title",
    "code_fragment": "Source code snippet",
    "explanation": "Detailed explanation text",
    "tags": ["tag1", "tag2", "tag3"]
}
```

Fields:

- **id**: Unique identifier (language-number format)
- **language**: Programming language (lowercase)
- **title**: Human-readable title describing code purpose
- **code_fragment**: Complete, runnable code snippet
- **explanation**: Educational text explaining code functionality, patterns, and concepts
- **tags**: List of keywords for categorization and retrieval

Appendix E: Code Excerpts

Language Detection Logic:

The LanguageDetector class implements priority-based matching:

```
def detect(self, code: str) -> str:
    lowered = code.lower()

    # Python detection (most distinct syntax)
    if any(indicator.lower() in lowered for indicator in
self.PYTHON_INDICATORS):
        if "class " in code and "{" in code and ":" in code:
```

```

        return "c++"
    return "python"

    # C++ specific features
    if any(indicator.lower() in lowered for indicator in self.CPP_INDICATORS):
        return "c++"

    # C-only features
    if any(indicator.lower() in lowered for indicator in self.C_INDICATORS):
        return "c"

    # Default fallback
    return "c"

```

Vector Search Implementation:

Cosine similarity computed via normalized dot product:

```

def search(self, query: str, top_k: int = 3) -> list[dict]:
    query_vec = self.model.encode([query], normalize_embeddings=True)[0]
    scores = np.dot(self.embeddings, query_vec)
    indices = np.argsort(scores)[::-1][:top_k]
    return [self.entries[idx] for idx in indices]

```

Appendix F: Testing Results Summary

Test Scenario	Input	Expected	Result	Status
Fibonacci	Recursive C	Detect C, explain recursion	Correct detection and explanation	PASS
Quicksort	Python impl	Detect Python, identify algorithm	Accurate language and pattern	PASS
Smart Pointer	C++ template	Detect C++, explain RAII	Correct language and concept	PASS
Minimal Code	Single printf	Default to C, provide explanation	Defaulted as designed	PASS
Pointer Arithmetic	C pointers	Retrieve pointer examples	Retrieved relevant references	PASS
Async Function	Python async/await	Explain concurrency	Correct async explanation	PASS
Ingestion	New example	Add to knowledge base	Successfully added	PASS
Empty Input	5 characters	Validation error	Appropriate error message	PASS

Performance Metrics:

- Average response time: 150ms
- Retrieval accuracy (covered patterns): 85%
- Language detection accuracy: 95%
- Knowledge base load time: 2.5s